### Red-Black tree

- Recall binary search tree
  - Key values in the left subtree <= the node value</p>
  - Key values in the right subtree >= the node value

#### Operations:

- insertion, deletion
- Search, maximum, minimum, successor, predecessor.
- O(h), h is the height of the tree.

#### Red-black trees

- Definition: a binary tree, satisfying:
  - 1. Every node is red or black
  - 2. The root is black
  - 3. Every leaf is NIL and is black
  - 4. If a node is red, then both its children are black
  - For each node, all paths from the node to descendant leaves contain the same number of black nodes.
- Purpose: keep the tree balanced.
- Other balanced search tree:
  - AVL tree, 2-3-4 tree, Splay tree, Treap

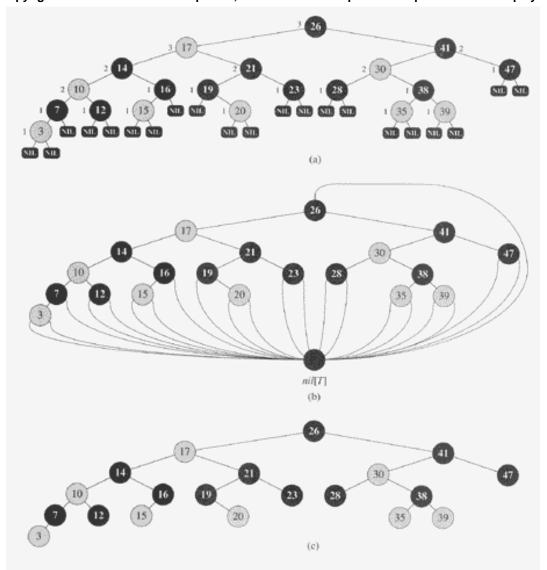


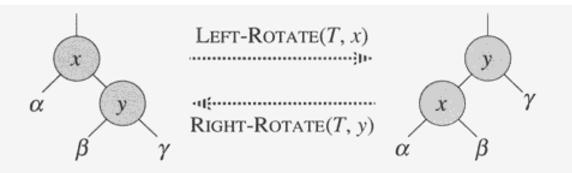
Figure 13.1 A red-black tree with black nodes darkened and red nodes shaded. Every node in a red-black tree is either red or black, the children of a red node are both black, and every simple path from a node to a descendant leaf contains the same number of black nodes. (a) Every leaf, shown as a NIL, is black. Each non-NIL node is marked with its black-height; NIL's have black-height 0. (b) The same red-black tree but with each NIL replaced by the single sentinel nil[T], which is always black, and with black-heights omitted. The root's parent is also the sentinel. (c) The same red-black tree but with leaves and the root's parent omitted entirely. We shall use this drawing style in the remainder of this chapter.

# Fields and property

- Left, right, ,parent, color, key
- bh(x), black-height of x, the number of black nodes on any path from x (excluding x) to a leaf.
- A red-black tree with n internal nodes has height at most 2log(n+1).
  - Note: internal nodes: all normal key-bearing nodes.
     External nodes: Nil nodes or the Nil Sentinel.
  - A subtree rooted at x contains at least 2<sup>bh(x)</sup>-1 internal nodes.
  - By property 4, bh(root)≥h/2.
  - $n \ge 2^{h/2}-1$

# Some operations in log(n)

- Search, minimum, maximum, successor, predecessor.
- Let us discuss insert or delete.



Right rotation: x=left[y]; left[y]=right[x]; If(right[x]!=nil) p[right[x]]=y; p[x]=p[y]; if(p[y]==nil) root=x; If(left[p[y]]=y) left[p[y]]=x; else right[p[y]]=x; right[x]=y; p[y]=x;

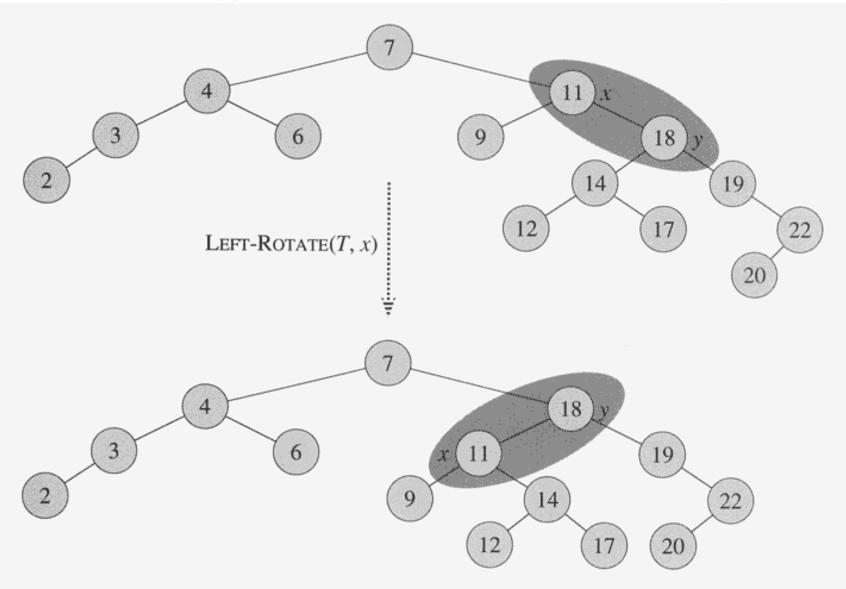
**Figure 13.2** The rotation operations on a binary search tree. The operation LEFT-ROTATE(T,x) transforms the configuration of the two nodes on the left into the configuration on the right by changing a constant number of pointers. The configuration on the right can be transformed into the configuration on the left by the inverse operation RIGHT-ROTATE(T,y). The letters  $\alpha$ ,  $\beta$ , and  $\gamma$  represent arbitrary subtrees. A rotation operation preserves the binary-search-tree property: the keys in  $\alpha$  precede key[x], which precedes the keys in  $\beta$ , which precede key[y], which precedes the keys in  $\gamma$ .

```
Left rotation:

y=right[x]; right[x] \leftarrow left[y]; lf(left[y]!=nil) p[left[y]]=x; p[y]=p[x];

if(p[x]==nil) {root=y;} else if (left[p[x]]==x) left[p[x]]=y; else right[p[x]]=y;

left[y]=x; p[x]=y;
```



**Figure 13.3** An example of how the procedure LEFT-ROTATE(T, x) modifies a binary search tree. Inorder tree walks of the input tree and the modified tree produce the same listing of key values.

Copyright © The McGraw-Hill Companies, Inc. Permission required for reproduction or display.

```
RB-INSERT(T, z)
    y \leftarrow nil[T]
 2 x \leftarrow root[T]
    while x \neq nil[T]
           do y \leftarrow x
                if key[z] < key[x]
 6
                  then x \leftarrow left[x]
                  else x \leftarrow right[x]
 8
     p[z] \leftarrow y
 9
     if y = nil[T]
10
         then root[T] \leftarrow z
11
         else if key[z] < key[y]
12
                  then left[y] \leftarrow z
13
                  else right[y] \leftarrow z
14
    left[z] \leftarrow nil[T]
15 right[z] \leftarrow nil[T]
16 color[z] \leftarrow RED
17
    RB-INSERT-FIXUP(T, z)
```

## Properties violations

- Property 1 (each node black or red): hold
- Proper 3: (each leaf is black sentinel): hold.
- Property 5: same number of blacks: hold
- Property 2: (root is black), not, if z is root (and colored red).
- Property 4: (the child of a red node must be black), not, if z's parent is red.

```
Case 1,2,3: p[z] is the left child of p[p[z]].
RB-INSERT-FIXUP(T, z)
                                             Correspondingly, there are 3 other cases,
     while color[p[z]] = RED
                                              In which p[z] is the right child of p[p[z]]
           do if p[z] = left[p[p[z]]]
                 then y \leftarrow right[p[p[z]]]
                       if color[y] = RED
 5
                         then color[p[z]] \leftarrow BLACK

    Case 1

 6
                               color[y] \leftarrow BLACK

    Case 1

                               color[p[p[z]]] \leftarrow RED

    Case 1

 8
                               z \leftarrow p[p[z]]

    Case 1

 9
                         else if z = right[p[z]]
10
                                  then z \leftarrow p[z]

    Case 2

11
                                        LEFT-ROTATE(T, z)

    Case 2

12
                               color[p[z]] \leftarrow BLACK

    Case 3

13
                               color[p[p[z]]] \leftarrow RED

    Case 3

14
                               RIGHT-ROTATE(T, p[p[z]])
                                                                                   Case 3
15
                 else (same as then clause
                               with "right" and "left" exchanged)
     color[root[T]] \leftarrow BLACK
```

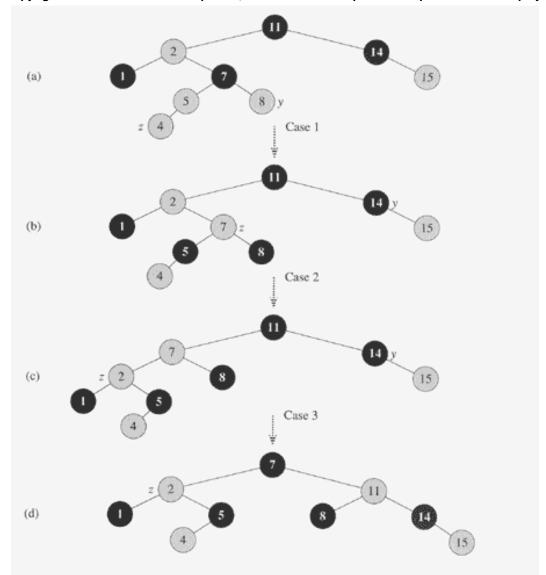
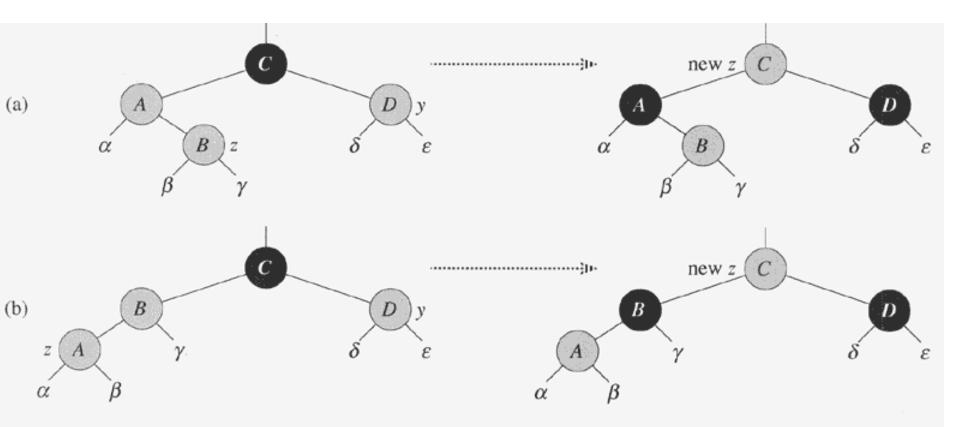
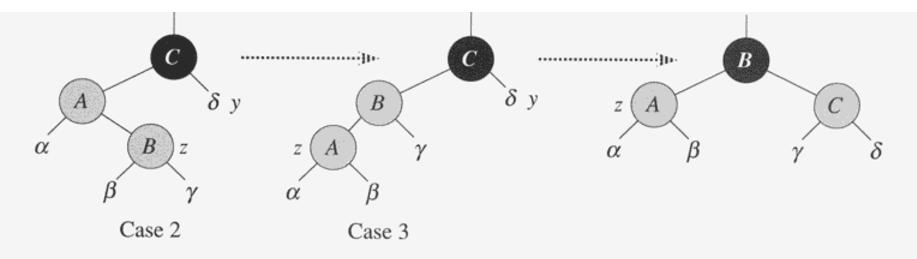


Figure 13.4 The operation of RB-INSERT-FIXUP. (a) A node z after insertion. Since z and its parent p[z] are both red, a violation of property 4 occurs. Since z's uncle y is red, case 1 in the code can be applied. Nodes are recolored and the pointer z is moved up the tree, resulting in the tree shown in (b). Once again, z and its parent are both red, but z's uncle y is black. Since z is the right child of p[z], case 2 can be applied. A left rotation is performed, and the tree that results is shown in (c). Now z is the left child of its parent, and case 3 can be applied. A right rotation yields the tree in (d), which is a legal red-black tree.



case 1: z's uncle is red.

Figure 13.5 Case 1 of the procedure RB-INSERT. Property 4 is violated, since z and its parent p[z] are both red. The same action is taken whether (a) z is a right child or (b) z is a left child. Each of the subtrees  $\alpha$ ,  $\beta$ ,  $\gamma$ ,  $\delta$ , and  $\varepsilon$  has a black root, and each has the same black-height. The code for case 1 changes the colors of some nodes, preserving property 5: all downward paths from a node to a leaf have the same number of blacks. The **while** loop continues with node z's grandparent p[p[z]] as the new z. Any violation of property 4 can now occur only between the new z, which is red, and its parent, if it is red as well.



Case 2: z's uncle is black and z is a right child. Case 3: z's uncle is black and z is a left child

**Figure 13.6** Cases 2 and 3 of the procedure RB-INSERT. As in case 1, property 4 is violated in either case 2 or case 3 because z and its parent p[z] are both red. Each of the subtrees  $\alpha$ ,  $\beta$ ,  $\gamma$ , and  $\delta$  has a black root ( $\alpha$ ,  $\beta$ , and  $\gamma$  from property 4, and  $\delta$  because otherwise we would be in case 1), and each has the same black-height. Case 2 is transformed into case 3 by a left rotation, which preserves property 5: all downward paths from a node to a leaf have the same number of blacks. Case 3 causes some color changes and a right rotation, which also preserve property 5. The **while** loop then terminates, because property 4 is satisfied: there are no longer two red nodes in a row.

What is the running time of RB\_INSERT\_FIX? And RB\_INSERT?

```
RB-DELETE(T, z)
     if left[z] = nil[T] or right[z] = nil[T]
        then y \leftarrow z
        else y \leftarrow \text{TREE-SUCCESSOR}(z)
    if left[y] \neq nil[T]
        then x \leftarrow left[y]
        else x \leftarrow right[y]
 7 p[x] \leftarrow p[y]
    if p[y] = nil[T]
        then root[T] \leftarrow x
10 else if y = left[p[y]]
11
                 then left[p[y]] \leftarrow x
12
                 else right[p[y]] \leftarrow x
13
     if y \neq z
        then key[z] \leftarrow key[y]
14
15
              copy y's satellite data into z
    if color[y] = BLACK
16
17
        then RB-DELETE-FIXUP(T, x)
18
     return y
```

```
RB-DELETE-FIXUP(T, x)
      while x \neq root[T] and color[x] = BLACK
           do if x = left[p[x]]
 3
                 then w \leftarrow right[p[x]]
 4
                       if color[w] = RED
 5
                          then color[w] \leftarrow BLACK

    Case 1

 6
                                color[p[x]] \leftarrow RED

    Case 1

                                LEFT-ROTATE (T, p[x])

    Case 1

 8
                                w \leftarrow right[p[x]]

    Case 1

 9
                       if color[left[w]] = BLACK and color[right[w]] = BLACK
10
                          then color[w] \leftarrow RED

    Case 2

11
                                x \leftarrow p[x]

    Case 2

12
                          else if color[right[w]] = BLACK
13
                                  then color[left[w]] \leftarrow BLACK

    Case 3

14
                                        color[w] \leftarrow RED

    Case 3

15
                                        RIGHT-ROTATE(T, w)

    Case 3

16
                                        w \leftarrow right[p[x]]
                                                                                   ⊳ Case 3
17
                                color[w] \leftarrow color[p[x]]

    Case 4

18
                                color[p[x]] \leftarrow BLACK

    Case 4

19
                                color[right[w]] \leftarrow BLACK

    Case 4

20
                                LEFT-ROTATE(T, p[x])

    Case 4

21
                               x \leftarrow root[T]

    Case 4

                 else (same as then clause with "right" and "left" exchanged)
22
23
     color[x] \leftarrow BLACK
```

#### Copyright © The McGraw-Hill Companies, Inc. Permission required for reproduction or display.

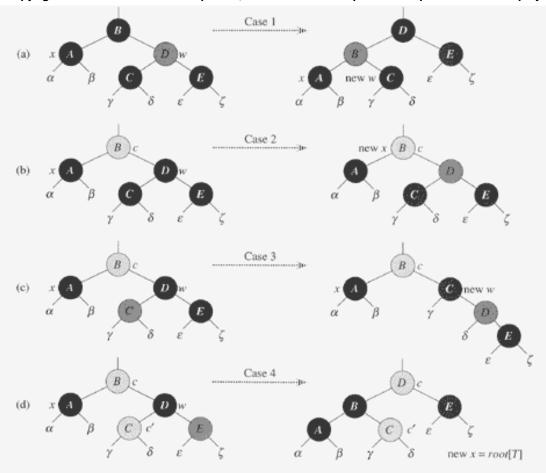


Figure 13.7 The cases in the white loop of the procedure RB-DELETE-FIXUP. Darkened nodes have *color* attributes BLACK, heavily shaded nodes have *color* attributes RED, and lightly shaded nodes have *color* attributes represented by c and c', which may be either RED or BLACK. The letters  $\alpha, \beta, \ldots, \zeta$  represent arbitrary subtrees. In each case, the configuration on the left is transformed into the configuration on the right by changing some colors and/or performing a rotation. Any node pointed to by x has an extra black and is either doubly black or red-and-black. The only case that causes the loop to repeat is case 2. (a) Case 1 is transformed to case 2, 3, or 4 by exchanging the colors of nodes B and D and performing a left rotation. (b) In case 2, the extra black represented by the pointer x is moved up the tree by coloring node D red and setting x to point to node B. If we enter case 2 through case 1, the while loop terminates because the new node x is red-and-black, and therefore the value c of its *color* attribute is RED. (c) Case 3 is transformed to case 4 by exchanging the colors of nodes C and D and performing a right rotation. (d) In case 4, the extra black represented by x can be removed by changing some colors and performing a left rotation (without violating the red-black properties), and the loop terminates.