


TRANSACTION MANAGEMENT

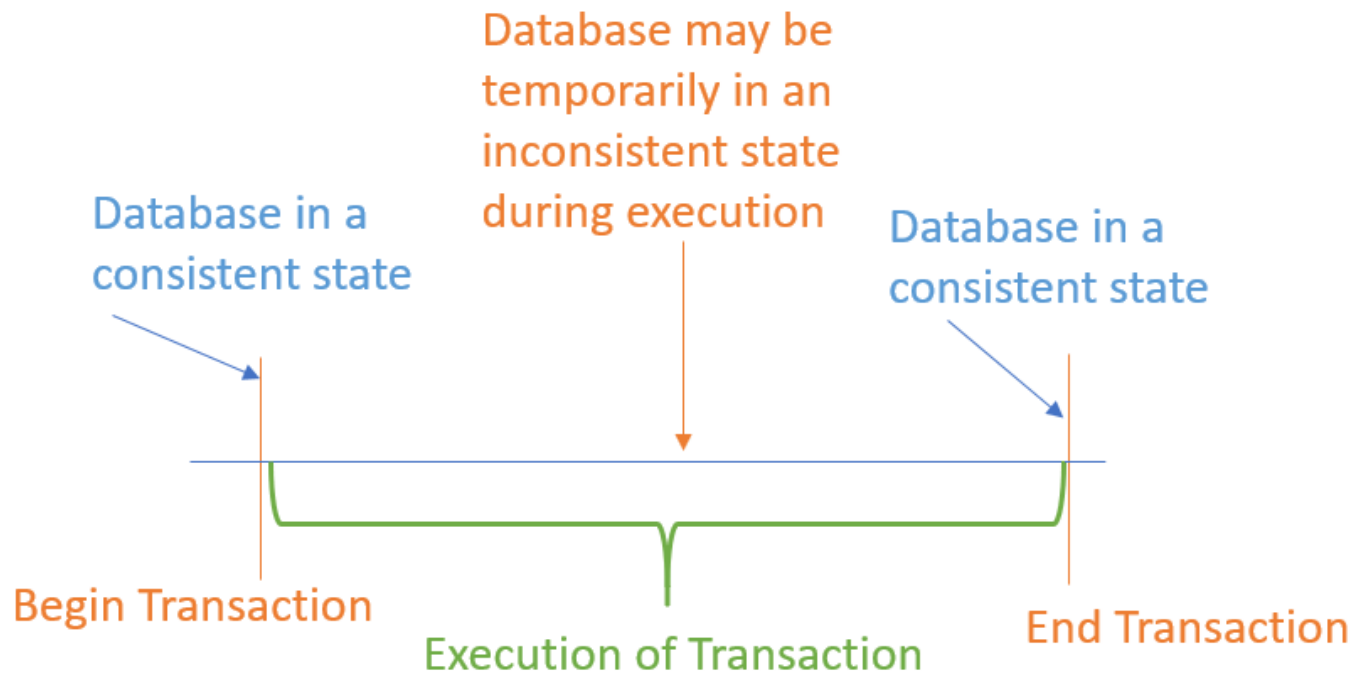


TRANSACTION CONCEPT

- **A transaction**
 - is a *unit* of program execution that accesses and possibly updates various data items.
 - must see a consistent database.
- Multiple transactions can execute in parallel.



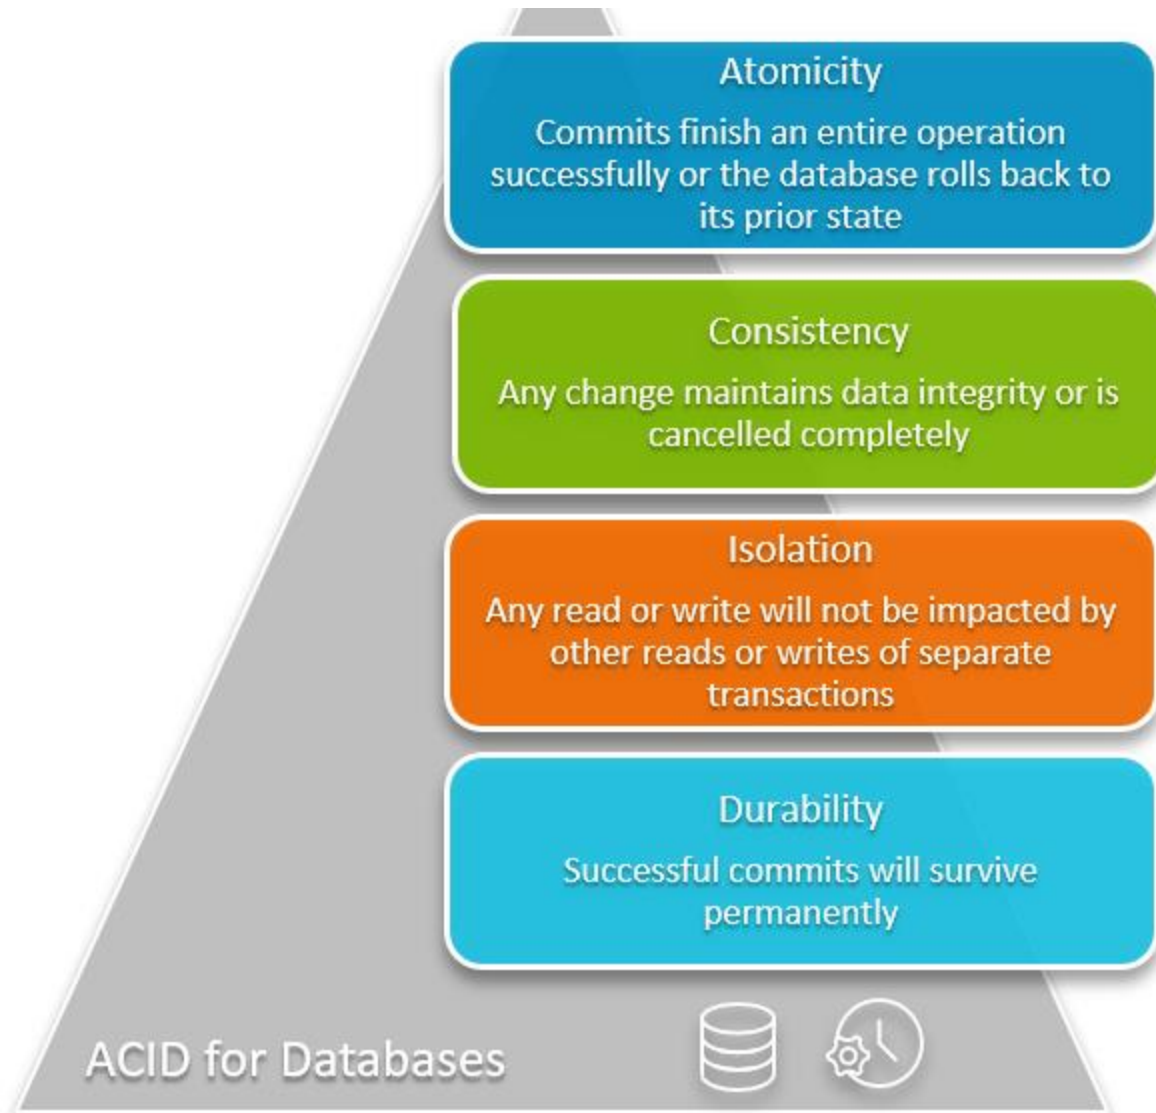
- **During** transaction execution the database may be temporarily inconsistent.
- When the transaction **completes** successfully (is committed), the database must be consistent.
- **After** a transaction commits, the changes it has made to the database persist, even if there are system failures.



- Two main issues to deal with:
 - Failures of various kinds, such as hardware failures and system crashes
 - Concurrent execution of multiple transactions



INTEGRITY PRESERVATION - TRANSACTION



ACID PROPERTIES

To preserve the integrity of data the database system must ensure:

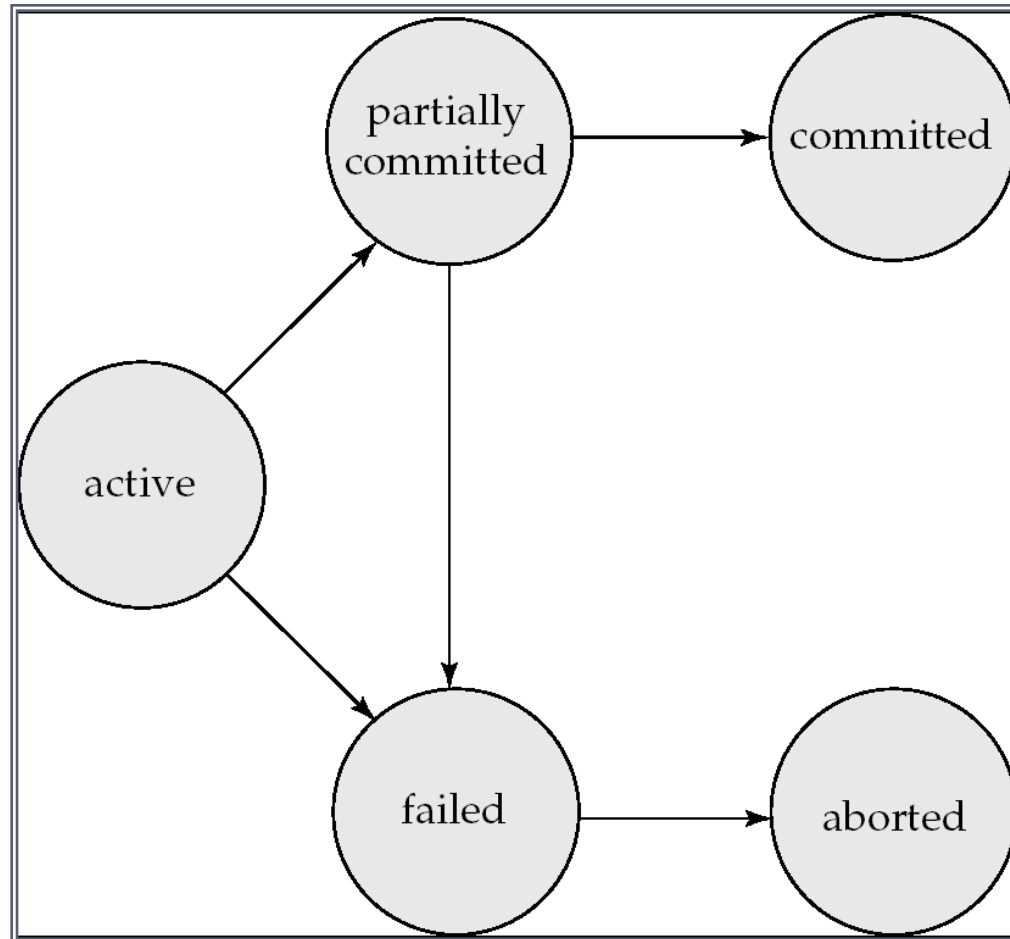
- **Atomicity.** Either all operations of the transaction are properly reflected in the database or none are.
- **Consistency.** Execution of a transaction run by itself preserves the consistency of the database.



- **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.
 - That is, for every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started, or T_j started execution after T_i finished.
- **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.



TRANSACTION STATES



TRANSACTION STATE

- **Active** – the initial state; the transaction stays in this state while it is executing
- **Partially committed** – after the final statement has been executed.
- **Failed** – after the discovery that normal execution can no longer proceed.
- **Aborted** – after the transaction has been rolled back and the database restored to its state prior to the start of the transaction. Two options after it has been aborted:
 - restart the transaction; can be done only if no internal logical error
 - kill the transaction
- **Committed** – after successful completion.



EXAMPLE OF FUND TRANSFER

- Transaction to transfer \$50 from account A to account B:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B)
 5. $B := B + 50$
 6. **write**(B)

- Atomicity requirement
- Consistency requirement
- Isolation requirement
- Durability requirement



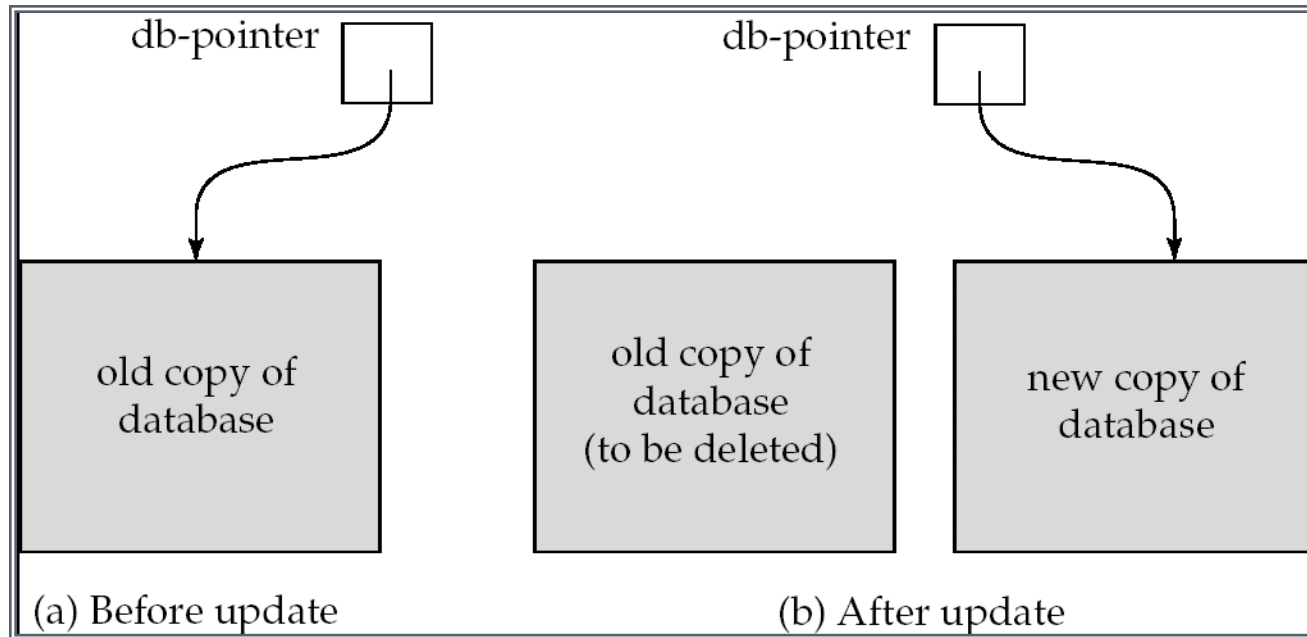
IMPLEMENTATION OF ATOMICITY AND DURABILITY

- The **recovery-management** component of a database system implements the support for atomicity and durability.
- The ***shadow-database*** scheme:
 - assume that only one transaction is active at a time.
 - a pointer called **db_pointer** always points to the current consistent copy of the database.
 - all updates are made on a *shadow copy* of the database, and **db_pointer** is made to point to the updated shadow copy only after the transaction reaches partial commit and all updated pages have been flushed to disk.
 - in case transaction fails, old consistent copy pointed to by **db_pointer** can be used, and the shadow copy can be deleted.



IMPLEMENTATION OF ATOMICITY AND DURABILITY (CONT.)

The shadow-database scheme:



- Assumes disks do not fail
- Useful for text editors, but
 - extremely inefficient for large databases
 - Does not handle concurrent transactions



CONCURRENT EXECUTIONS

- Multiple transactions are allowed to run concurrently in the system.
- Advantages are:
 - **increased processor and disk utilization**, leading to better transaction *throughput*
 - **reduced average response time** for transactions



SCHEDULES

- **Schedule** – a sequences of instructions that specify the order of concurrent transactions are executed
- COMMIT
- FAIL / ABORT



SCHEDULE 1

- Let T_1 transfer \$50 from A to B , and T_2 transfer 10% of the balance from A to B .
- A serial schedule in which T_1 is followed by T_2 :

T_1	T_2
read(A) $A := A - 50$ write (A) read(B) $B := B + 50$ write(B)	read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B)



SCHEDULE 2

- A serial schedule where T_2 is followed by T_1

T_1	T_2
read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B)	read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B)



SCHEDULE 3

- Let T_1 and T_2 be the transactions defined previously. The following schedule is not a serial schedule, but it is *equivalent* to Schedule 1.

T_1	T_2
read(A) $A := A - 50$ write(A)	read(A) $temp := A * 0.1$ $A := A - temp$ write(A)
read(B) $B := B + 50$ write(B)	read(B) $B := B + temp$ write(B)

In Schedules 1, 2 and 3, the sum $A + B$ is preserved.



SCHEDULE 4

- The following concurrent schedule does not preserve the value of $(A + B)$.

T_1	T_2
read(A) $A := A - 50$	read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B)
write(A) read(B) $B := B + 50$ write(B)	$B := B + temp$ write(B)



CONCURRENT EXECUTIONS OF TRANSACTIONS AND RELATED PROBLEMS

- Dirty Read Problem
- Lost Update Problem
- Unrepeatable Read Problem
- Phantom Read Problem
- Incorrect Aggregate Problem



W R CONFLICTS

DIRTY READ !!!

- Reading Uncommitted Data
- Temporary Update Problem

T1	T2
R(A)	
A=A-2;	
W(A)	
	R(A)
	A=A-5
	W(A)
	Commit
R(A)	
W(A)	



WW CONFLICT

LOST UPDATE – BLIND WRITE

T1	T2
R(A)	
A=A+5;	
W(A)	
	A=A+10
	W(A)
	Commit
Commit	



R W CONFLICTS

UNREPEATABLE READS

T1	T2
R(A)	
	R(A)
A=A-5	
W(A)	
Commit	
	R(A)



CONFLICTS

PHANTOM READ

T1	T2
R(A)	
	R(A)
Delete(A)	
	R(A)



CONFLICTS

INCORRECT SUMMARY

T1	T2
<pre>read_item(X) X = X - N write_item(X)</pre> <pre>read_item(Y) Y = Y + N write_item(Y)</pre>	<pre>sum = 0 read_item(A) sum = sum + A</pre> <pre>read_item(X) sum = sum + X read_item(Y) sum = sum + Y</pre>



ANOMALIES WITH INTERLEAVED EXECUTION

- $R(A) \ R(A)$
- $W(A) \ R(A)$
- $R(A) \ W(A)$
- $W(A) \ W(A)$



- They belong to different transactions
- They operate on the same data item
- At least one of them is a write operation

CONFLICTING PAIRS !!!

Conflict Pairs
$W_1(A) \ R_2(A)$
$R_1(A) \ W_2(A)$
$W_1(A) \ W_2(A)$

Non – Conflict Pairs
$R(A) \ R(A)$
$R(B) \ R(A)$
$W(B) \ R(A)$
$R(B) \ W(A)$
$W(A) \ R(B)$



SWAP NON -CONFLICT PAIRS

ARE S AND S¹ CONFLICT EQUIVALENT?

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	

T1	T2
R(A)	
W(A)	
R(B)	
	R(A)
	W(A)



SERIALIZABILITY

- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule.

1. **Conflict Serializability**

2. **View Serializability**



CONFLICT SERIALIZABILITY

- If a schedule S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are **conflict equivalent**.
- We say that a schedule S is **conflict serializable** if it is conflict equivalent to a serial schedule



SWAP NON -CONFLICT PAIRS

CONFLICT SERIALIZABILITY

- Check if conflict serializable?

T_1	T_2
read(A) write(A)	read(A) write(A)
read(B) write(B)	
	read(B) write(B)

T_1	T_2
read(A) write(A) read(B) write(B)	read(A) write(A) read(B) write(B)



SWAP NON -CONFLICT PAIRS

CONFLICT SERIALIZABILITY


Check if conflict serializable?

T_3	T_4
read(Q)	write(Q)
write(Q)	



TESTING FOR SERIALIZABILITY

- Consider some schedule of a set of transactions T_1, T_2, \dots, T_n
- **Precedence graph** \rightarrow a direct graph where the vertices are the transactions
- Draw an arc from T_i to T_j if the two transaction conflict, and T_i accessed the data item on which the conflict arose earlier.
- Label the arc by the item that was accessed.



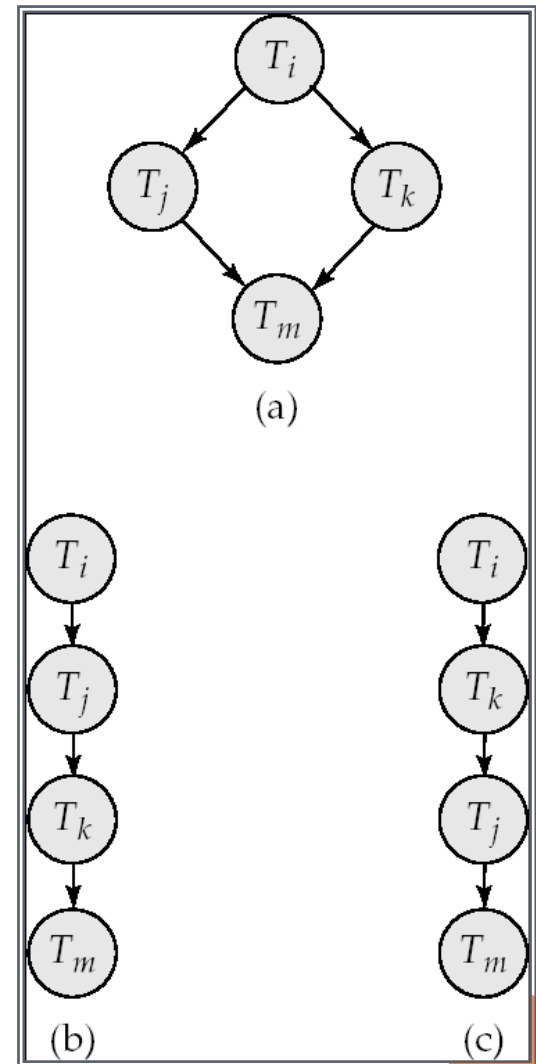
The set of edges consists of all edges

$T_i \rightarrow T_j$ for which one of three conditions holds:

1. T_i executes $write(Q)$ before T_j executes $read(Q)$.
2. T_i executes $read(Q)$ before T_j executes $write(Q)$.
3. T_i executes $write(Q)$ before T_j executes $write(Q)$.

TEST FOR CONFLICT SERIALIZABILITY

- A schedule is conflict serializable if and only if its precedence graph is **acyclic**.
- If precedence graph is acyclic, the serializability order can be obtained by a *topological sorting* of the graph.



Draw edges b/w conflict pairs

TEST FOR CONFLICT SERIALIZABILITY !

T1	T2	T3
R(x)		
		R(y)
		R(x)
	R(y)	
	R(z)	
		W(y)
	W(z)	
R(z)		
W(x)		
W(z)		



Draw edges b/w conflict pairs

TEST FOR CONFLICT SERIALIZABILITY !

T4	T5	T6
R(x)		
	W(x)	
W(x)		
		W(x)



VIEW SERIALIZABILITY

- A schedule S is **view serializable** if it is view equivalent to a serial schedule.
- Every conflict serializable schedule is also view serializable.
- Every view serializable schedule that is not conflict serializable has **blind writes**.



LETS CHECK !

- Below is a schedule which is view-serializable but *not* conflict serializable.

T4	T5	T6
R(x)		
	W(x)	
W(x)		
		W(x)



TESTING FOR VIEW SERIALIZABILITY

○ S and S' are **view equivalent** if the following three conditions are met:

- 1. Initial Read :** For each data item Q , if transaction T_i reads the initial value of Q in schedule S , then transaction T_i must, in schedule S' , also read the initial value of Q .
- 2. Intermediate Read:** For each data item Q if transaction T_i executes **read**(Q) in schedule S , and that value was produced by transaction T_j (if any), then transaction T_i must in schedule S' also read the value of Q that was produced by transaction T_j .
- 3. Final Write :** For each data item Q , the transaction (if any) that performs the final **write**(Q) operation in schedule S must perform the final **write**(Q) operation in schedule S' .

View equivalence is also based purely on **reads and writes** alone.

Initial Read
Intermediate Read
Final Write

TEST FOR VIEW SERIALIZABILITY !

T1	T2
R(a)	
W(a)	
	R(a)
	W(a)
R(b)	
W(b)	
	R(b)
	W(b)

POSSIBILITIES ???



Initial Read
Intermediate Read
Final Write

TEST FOR VIEW SERIALIZABILITY !

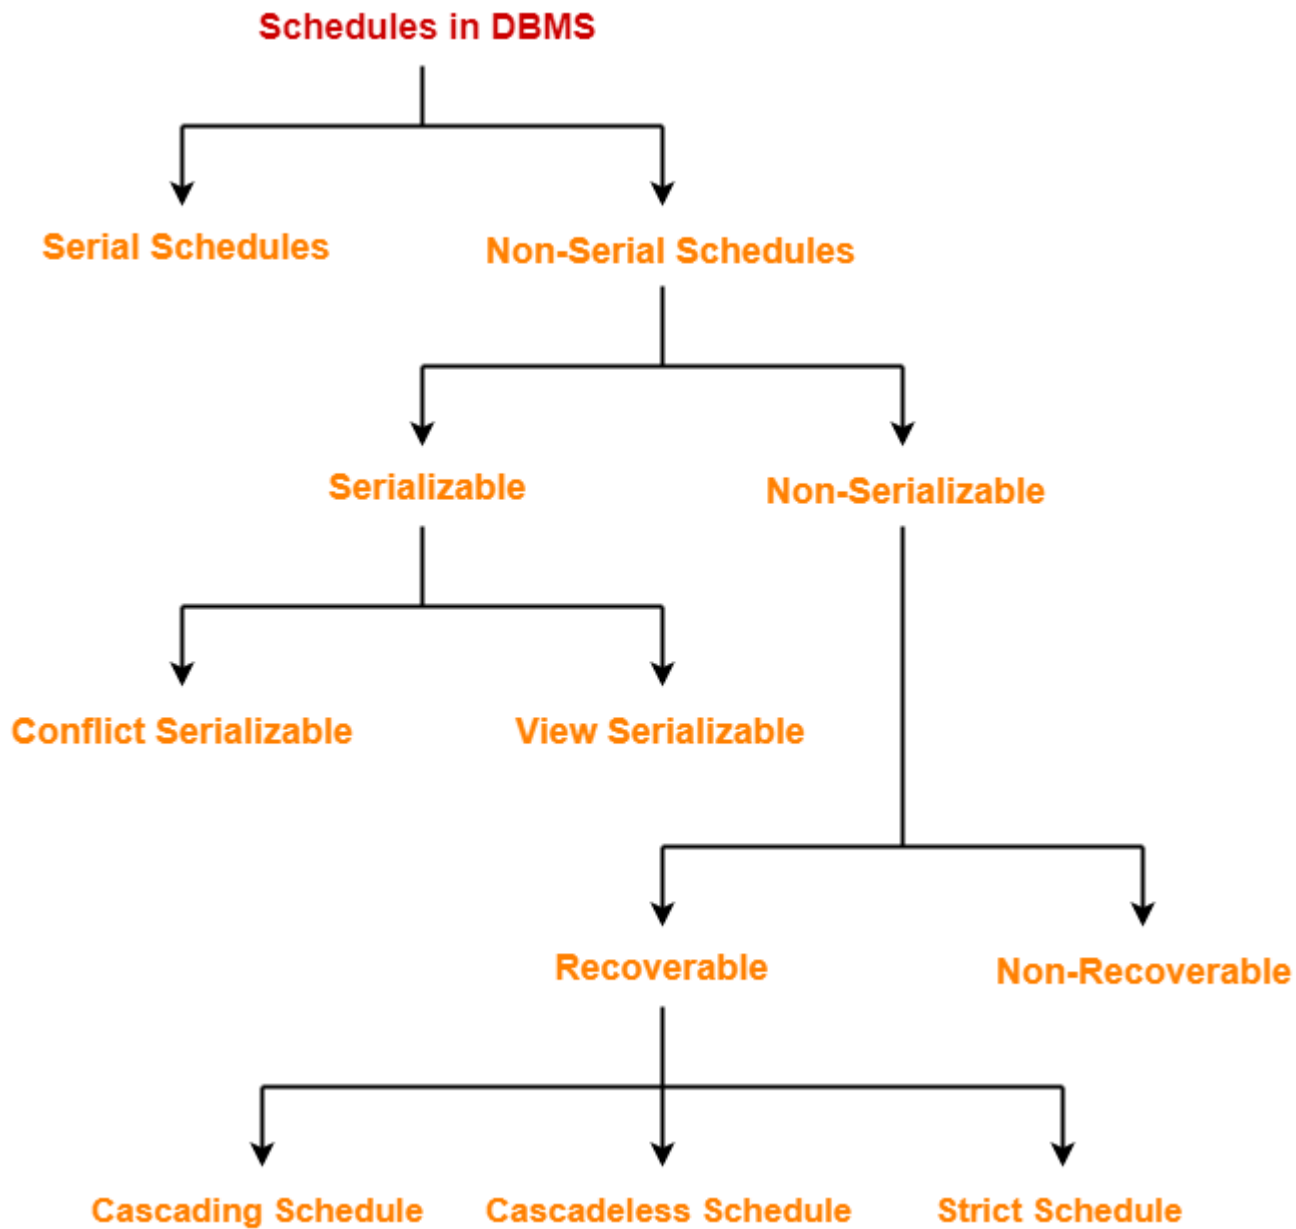
S

T1	T2
R(a)	
W(a)	
	R(a)
	W(a)
R(b)	
W(b)	
	R(b)
	W(b)

S¹

T1	T2
R(a)	
W(a)	
R(b)	
W(b)	
	R(a)
	W(a)
	R(b)
	W(b)





SYSREM RECOVERY

- Failures
 - Transaction Failures
 - Logical error
 - System error
 - System Crash



NON- RECOVERABLE VS RECOVERABLE SCHEDULES

P

T1	T2
R(a)	
A=a-10	
W(a)	
	R(a)
	A=a-20
	W(a)

Q

T1	T2
R(a)	
A=a-10	
W(a)	
Commit	
	R(a)
	A=a-20
	W(a)



RECOVERABLE SCHEDULES

- **Recoverable schedule** — if a transaction T_j reads a data item previously written by a transaction T_i , then the commit operation of T_i appears before the commit operation of T_j .



CASCADING VS CASCADELESS SCHEDULE

T1	T2	T3
R(a)		
a=a+10		
W(a)		
	R(a)	
	a=a+10	
	W(a)	
		R(a)
		a=a+10
		W(a)



CASCADING VS **CASCADELESS** SCHEDULE

T1	T2	T3
R(a)		
a=a+10		
W(a)		
Commit		
	R(a)	
	a=a+10	
	W(a)	
	Commit	
		R(a)
		a=a+10
		W(a)
		Commit



CASCADING ROLLBACKS

- **Cascading rollback** – a single transaction failure leads to a series of transaction rollbacks



CASCADELESS SCHEDULES

- **Cascadeless schedules** — cascading rollbacks cannot occur; for each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the read operation of T_j .
- Every cascadeless schedule is also recoverable
- It is desirable to restrict the schedules to those that are cascadeless



LOG BASED RECOVERY

- **The facility to maintain or recover data if any failure may occur in the system**
- **Log == Record**
- **Deferred Database Modification**
- **Immediate Database Modification**



FIELDS IN A LOG

- **Transaction identifier**
- **Data item**
- **Old value**
- **New value**

- **Types of logs**
 - **<Ti start>**
 - **<Ti commit>**
 - **<Ti abort>**



- New values are stored in the log
- Redo / No undo

DEFERRED DATABASE MODIFICATION / UPDATE

T1
R(A)
A=A+20
W(A)
R(B)
B=B+50
W(B)
Commit

T1
R(A)
A=A+20
W(A)
R(B)
B=B+50
W(B)



- New and old values are stored in the log
- Redo / Undo

IMMEDIATE DATABASE MODIFICATION / UPDATE

T1
R(A)
A=A+20
W(A)
R(B)
B=B+50
W(B)
Commit

T1
R(A)
A=A+20
W(A)
R(B)
B=B+50
W(B)



LOCK-BASED PROTOCOLS

- A lock is a mechanism to control concurrent access to a data item
 - 1. *exclusive (X) mode*. Data item can be both read as well as written. X-lock is requested using **lock-X** instruction.
 - 2. *shared (S) mode*. Data item can only be read. S-lock is requested using **lock-S** instruction.



SHARED (S) – READ
EXCLUSIVE(X) – READ & WRITE

T1
R(A)

T2
R(A)
A=A-50
W(A)



LOCK-BASED PROTOCOLS

- Lock-compatibility matrix

	S	X
S	true	false
X	false	false

T1	T2
R(A)	
	R(A)
	W(A)

T1	T2
W(A)	
	R(A)
	W(A)

T1	T2
W(A)	
	W(A)

INTENT LOCKING

- IS – Intent to get S lock(s) at finer granularity.
- IX – Intent to get X lock(s) at finer granularity.
- SIX mode: Like S & IX at the same time.



PITFALLS IN LOCK BASED PROTOCOLS

- Serialization **may not** be guaranteed
- **May not** be free from non - recoverability
- **May not** be free from Deadlocks
- **May not** be free from Starvation



SERIALIZATION MAY NOT BE GUARANTEED

T1	T2
R(A) W(A)	
	R(A)
R(B) W(B)	




MAY NOT BE FREE FROM NON- RECOVERABILITY

T1	T2
R(A) W(A)	
	R(A)
	...
	...
*	



MAY NOT BE FREE FROM DEADLOCKS

- Waiting indefinitely for resources



MAY NOT BE FREE FROM STARVATION

T1	T2	T3	T4
X(A)	S(A)		



THE TWO-PHASE LOCKING PROTOCOL

- This is a protocol which ensures conflict-serializable schedules.
- Phase 1: Growing Phase
 - transaction may
 - obtain locks
 - not release locks
- Phase 2: Shrinking Phase
 - transaction may
 - release locks
 - not obtain locks



EXAMPLE

T2

R(A)

W(A)

R(B)

R(C)

W(D)

R(D)



DEMONSTRATION OF 2PL WITH LOCKPOINTS

T1	T2
S(A)	
	S(A)
X(B)	
U(A)	
	X(D)
U(B)	
	U(A)
	U(D)



TYPES OF 2PL

- Strict 2-PL
- Rigorous 2-PL
- Conservative 2-PL



STRICT 2 PL

- **All Exclusive(X) locks** held by the transaction can be released until *after* the Transaction Commits.
- Helps in
 - Recoverability
 - Cascadeless



RIGOROUS 2 PL

- **All Exclusive(X) and Shared(S) locks** held by the transaction can be released until *after* the Transaction Commits.



CONSERVATIVE 2PL

- Lock all items it needs then transaction starts execution
 - If any locks can not be obtained, then do not lock anything
- Difficult but deadlock free



- Transaction concepts, properties of transactions, serializability of transactions, testing for serializability, System recovery, Two-Phase Commit protocol, Recovery and Atomicity, Log-based recovery, concurrent executions of transactions and related problems, Locking mechanism, solution to concurrency related problems, deadlock, , two-phase locking protocol, Isolation, Intent locking

