

Heuristic Search

- INTRODUCTION
- HILL-CLIMBING
- DYNAMIC PROGRAMMING
- THE BEST-FIRST SEARCH ALGORITHM
- ADMISSIBILITY, MONOTONICITY, AND INFORMEDNESS
- USING HEURISTICS IN GAMES.

Heuristic

- Heuristic – the study of the methods and rules of discovery and invention
- State Space Heuristics – Formalized as rules for choosing those branches in a state space that are most likely to lead to an acceptable problem solution
- Apply Heuristics When:
 - A problem is ambiguous and may not have an EXACT solution
 - Eg: Medical diagnosis
 - The computational cost of finding an exact solution is prohibitive.
 - Eg: Chess – where the number of possible states increases exponentially or factorially with the depth of the search.
- Heuristics guide the search along the most “promising” path through space – Informed Search.
- Eliminates unpromising states and their descendants from consideration and finds an acceptable solution



Inherent limitation of heuristic search

- Heuristic is only an informed guess of the next step to be taken in solving a problem
- Heuristics use limited information, so they are seldom able to predict the exact behavior of the state space farther along in the search
- A heuristic can lead a search algorithm to a suboptimal solution or fail to find any solution at all

Importance of heuristics


- It is not feasible to examine every inference that can be made in a mathematics domain
- Heuristic search is often the only practical answer
- Reduce complex information to a simple and manageable set of choices
- Help people turn an intention into a realized action
- Provide quick and relatively inexpensive feedback to designers

TIC-TAC-TOE




- Total number of states that need to be considered in an exhaustive search at $9 \times 8 \times 7 \times \dots$ or $9!$
- Symmetry reductions on the second level further reduce the number of paths through the space to $12 \times 7!$


The “most wins” heuristic applied to the first children in tic-tac-toe.



Three wins through
a corner square




Four wins through
the center square




Two wins through
a side square

Heuristically reduced state space for tic-tac-toe.



Hill climbing

- Hill climbing strategies expand the current state in the search and evaluate its children
- The best child is selected for further expansion; neither its siblings nor its parent are retained
- Search halts when it reaches a state that is better than any of its children
- Go uphill along the steepest possible path until it can go no farther
- Keeps no history, hence the algorithm cannot recover from failures of its strategy




Hill-climbing search: 8-Queens problem

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♚	13	16	13	16
♚	14	17	15	♚	14	16	16
17	♚	16	18	15	♚	15	♚
18	14	♚	15	15	14	♚	16
14	14	13	17	12	14	12	18

- h = number of pairs of queens that are attacking each other, either directly or indirectly
- $h = 17$ for the above state

Hill-climbing search: 8-Queens problem



- A local minimum with $h = 1$

Dynamic Programming

- Sometimes called the forward-backward. or, when using probabilities, the Viterbi algorithm.
- DP keeps track of and reuses subproblems already searched and solved within the solution of the larger problem.
- Used in
 - Optimal Global Alignment
 - Minimum Edit Distance Between two strings

Optimal Global Alignment

■ Small Example

String #1

- BAADDCABDDA

String #2

- BBADCBA

■ Rules

Cannot change order of respective elements

Can have spaces between elements

■ Possible solutions

How do we figure out optimal solution?


OPTIMAL GLOBAL ALIGNMENT

	—	B	A	A	D	D	C	A	B	D	D	A
—	0	1	2	3	4	5	6	7	8	9	10	11
B	1	0										
B	2											
A	3											
D	4											
C	5											
B	6											
A	7											

The forward stage

- Fill the array from the upper left corner
- The value of (x,y) is a function
 - $\min \text{cost}((x-1,y), (x-1,y-1), (x, y-1))$
- If there is a match: Add 0 to $(x-1,y-1)$
- If there is no match: Add 2 to $(x-1,y-1)$
- If we shift: Add 1 to the previous column
- If we insert a character: Add 1 to the previous row

	—	B	A	A	D	D	C	A	B	D	D	A
—	0	1	2	3	4	5	6	7	8	9	10	11
B	1	0	1	2	3	4	5	6	7	8	9	10
B	2	1	2	3	4	5	6	7	6	7	8	9
A	3	2	1	2	3	4	5	6	7	8	9	8
D	4	3	2	3	2	3	4	5	6	7	8	9
C	5	4	3	4	3	4	3	4	5	6	7	8
B	6	5	6	5	4	5	4	5	4	5	6	7
A	7	6	5	4	5	6	5	4	5	6	7	6



The backward stage

- Once the matrix is filled
- From the best alignment count, we produce a specific alignment of characters
- Begin at the lower right-hand corner
- Move back through the matrix
- Each step, select one of the immediate state's predecessors (previous diagonal, row, or column) → Choose the minimum

BAADDCABDDA
BBADC B A

Min Edit Distance Spell Checker

- What words from our dictionary best approximate a word we do not recognize (misspelled word)
- We need to know the “distance” between two words Minimum edit distance
- The number of insertions, deletions and replacements to turn the **source word** into the **target word**

	—	e	x	e	c	u	t	i	o	n
—	0	1	2	3	4	5	6	7	8	9
i	1	2	3	4	5	6	7	8	9	10
n	2	3	4	5	6	7	8	9	10	11
t	3	4	5	6	7	8	9	10	11	12
e	4	5	6	5	6	7	8	9	10	11
n	5	6	7	6	7	8	9	10	11	12
t	6	7	8	7	8	9	8	9	10	11
i	7	8	9	8	9	10	9	8	9	10
o	8	9	10	9	10	11	10	9	8	9
n	9	10	11	10	11	12	11	10	9	8

cost of (x,y) is the minimum of

- .Cost of (x-1,y) + insertion
- .Cost of (x-1,y-1) + replacement
- .Cost of (x, y-1) + deletion

intention – source word

execution – target word

Our “cost”

1 for a character insertion or deletion

2 for a replacement (deletion + insertion)

intention

ntention

delete i, cost 1

etention

replace n with e, cost 2

exention

replace t with x, cost 2

exenution

insert u, cost 1


execution

replace n with c, cost 2

Best-First Search Algorithm

- Hill climbing tends to become stuck at local maxima
- If they reach a state that has a better evaluation than any of its children, the algorithm halts
- Hill climbing can be used effectively if the evaluation function is sufficiently informative to avoid local maxima and infinite paths
- Heuristic search requires a more flexible algorithm: this is provided by best-first search, where, with a priority queue, recovery from local maxima is possible

Heuristic search of a hypothetical state space.



1. **open = [A5]; closed = []**
2. **evaluate A5; open = [B4,C4,D6]; closed = [A5]**
3. **evaluate B4; open = [C4,E5,F5,D6]; closed = [B4,A5]**
4. **evaluate C4; open = [H3,G4,E5,F5,D6]; closed = [C4,B4,A5]**
5. **evaluate H3; open = [O2,P3,G4,E5,F5,D6]; closed = [H3,C4,B4,A5]**
6. **evaluate O2; open = [P3,G4,E5,F5,D6]; closed = [O2,H3,C4,B4,A5]**
7. **evaluate P3; the solution is found!**

```


function best_first_search;

begin
  open := [Start];
  closed := [];
  while open ≠ [] do
    begin
      remove the leftmost state from open, call it X;
      if X = goal then return the path from Start to X
      else begin
        generate children of X;
        for each child of X do
          case
            the child is not on open or closed:
              begin
                assign the child a heuristic value;
                add the child to open
              end;
            the child is already on open:
              if the child was reached by a shorter path
              then give the state on open the shorter path
            the child is already on closed:
              if the child was reached by a shorter path then
              begin
                remove the state from closed;
                add the child to open
              end;
          end;
        put X on closed;
        re-order states on open by heuristic merit (best leftmost)
      end;
    end;
  return FAIL

```

- The best-first search algorithm always select the most promising state on open for further expansion
- it is using a heuristic that may prove erroneous, it does not abandon all the other states but maintains them on open
- In the event a heuristic leads the search down a path that proves incorrect, the algorithm will eventually retrieve some previously generated, “next best” state from open and shift its focus to another part of the space
- In best-first search, as in depth-first and breadth-first search algorithms, the open list allows backtracking from paths that fail to produce a goal

The start state, first set of moves, and goal state for an 8-puzzle instance.



Informed (Heuristic) Search Strategies

- **Informed Search** – a strategy that uses problem-specific knowledge beyond the definition of the problem itself
- **Best-First Search** – an algorithm in which a node is selected for expansion based on an evaluation function $f(n)$
 - Traditionally the node with the lowest evaluation function is selected
 - Choose the node that *appears* to be the best

Best-first search

- Idea: use an **evaluation function** $f(n)$ for each node
 - estimate of "desirability"
 - Expand most desirable unexpanded node
- Implementation:

Order the nodes in decreasing order of desirability
- Special cases:
 - greedy best-first search
 - A^* search

Greedy best-first search

- Evaluation function $f(n) = h(n)$ (**h**euristic)
= estimate of cost from n to *goal*
- e.g., $h_{SLD}(n)$ = straight-line distance from n to goal
- Greedy best-first search expands the node that **appears** to be closest to goal

A* search

- Idea: avoid expanding paths that are already expensive
- Evaluation function $f(n) = g(n) + h(n)$
 - $g(n)$ = cost so far to reach n
 - $h(n)$ = estimated cost from n to goal
 - $f(n)$ = estimated total cost of path through n to goal

Three heuristics applied to states in the 8-puzzle.

<table><tr><td>2</td><td>8</td><td>3</td></tr><tr><td>1</td><td>6</td><td>4</td></tr><tr><td></td><td>7</td><td>5</td></tr></table>	2	8	3	1	6	4		7	5	5	6	0
2	8	3										
1	6	4										
	7	5										
<table><tr><td>2</td><td>8</td><td>3</td></tr><tr><td>1</td><td></td><td>4</td></tr><tr><td>7</td><td>6</td><td>5</td></tr></table>	2	8	3	1		4	7	6	5	3	4	0
2	8	3										
1		4										
7	6	5										
<table><tr><td>2</td><td>8</td><td>3</td></tr><tr><td>1</td><td>6</td><td>4</td></tr><tr><td>7</td><td>5</td><td></td></tr></table>	2	8	3	1	6	4	7	5		5	6	0
2	8	3										
1	6	4										
7	5											
	Tiles out of place	Sum of distances out of place	2 x the number of direct tile reversals									

1	2	3
8		4
7	6	5

Goal

The heuristic f applied to states in the 8-puzzle.

$$g(n) = 0$$

Start

2	8	3
1	6	4
7		5

$$g(n) = 1$$

2	8	3
1	6	4
	7	5

2	8	3
1		4
7	6	5

2	8	3
1	6	4
7	5	

Values of $f(n)$ for each state,

6

4

6

where:


$$f(n) = g(n) + h(n),$$

$g(n)$ = actual distance from n
to the start state, and

$h(n)$ = number of tiles out of place.

1	2	3
8		4
7	6	5

Goal




Admissibility, Monotonicity & Informedness

Admissible heuristics

- A heuristic $h(n)$ is **admissible** if for every node n ,
 $h(n) \leq h^*(n)$, where $h^*(n)$ is the **true** cost to reach the goal state from n .
- An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is **optimistic**
- **Theorem:** If $h(n)$ is admissible, A^* using TREE-SEARCH is optimal

Optimality of A^*


- Suppose some suboptimal goal G_2 has been generated and is in the fringe. Let n be an unexpanded node in the fringe such that n is on a shortest path to an optimal goal G .



- $f(G_2) = g(G_2)$ since $h(G_2) = 0$
- $g(G_2) > g(G)$ since G_2 is suboptimal
- $f(G) = g(G)$ since $h(G) = 0$
- $f(G_2) > f(G)$ from above

Optimality of A^*

- Suppose some suboptimal goal G_2 has been generated and is in the fringe. Let n be an unexpanded node in the fringe such that n is on a shortest path to an optimal goal G .



- $f(G_2) > f(G)$ from above
- $h(n) \leq h^*(n)$ since h is admissible
- $g(n) + h(n) \leq g(n) + h^*(n)$
- $f(n) \leq f(G)$

Hence $f(G_2) > f(n)$, and A^* will never select G_2 for expansion

Consistent heuristics / Monotonicity


- A heuristic is **consistent/monotone** if for every node n , every successor n' of n generated by any action a ,

$$h(n) \leq c(n,a,n') + h(n')$$

- If h is consistent, we have

$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n,a,n') + h(n') \\ &\geq g(n) + h(n) \\ &= f(n) \end{aligned}$$

- i.e., $f(n)$ is non-decreasing along any path.
- **Theorem:** If $h(n)$ is consistent, A* using GRAPH-SEARCH is optimal



Admissible heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance
(i.e., no. of squares from desired location of each tile)

- $\underline{h_1(S)} = ?$

- $\underline{h_2(S)} = ?$

7	2	4
5		6
8	3	1

Start State


	1	2
3	4	5
6	7	8

Goal State

Admissible heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance
(i.e., no. of squares from desired location of each tile)



- $h_1(S) = ?$ 8
- $h_2(S) = ?$ $3+1+2+2+2+3+3+2 = 18$

Dominance/ Informedness


- If $h_2(n) \geq h_1(n)$ for all n (both admissible),
then h_2 **dominates** $h_1 \rightarrow h_2$ is better for search
- Typical search costs (average number of nodes expanded):
 - $d=12$, IDS = 3,644,035 nodes
 $A^*(h_1) = 227$ nodes
 $A^*(h_2) = 73$ nodes
 - $d=24$ IDS = too many nodes
 $A^*(h_1) = 39,135$ nodes
 $A^*(h_2) = 1,641$ nodes

Adversarial Search

Games vs. search problems


- "Unpredictable" opponent → specifying a move for every possible opponent reply
- Time limits → unlikely to find goal, must approximate

Game tree (2-player, deterministic, turns, zero-sum)



Minimax

- Perfect play for deterministic games
- Idea: choose move to position with highest **minimax value**
= best achievable payoff against best play
- E.g., 2-ply game:
-



Minimax algorithm

function MINIMAX-DECISION(*state*) *returns an action*

$v \leftarrow \text{MAX-VALUE}(\textit{state})$

return the *action* in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return *v*

function MIN-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)


$v \leftarrow \infty$

for *a, s* in SUCCESSORS(*state*) **do**


$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return *v*


α - β pruning example




α - β pruning example




α - β pruning example



α - β pruning example



α - β pruning example




Properties of α - β

- Pruning **does not** affect final result
- Good move ordering improves effectiveness of pruning
- With "perfect ordering," time complexity = $O(b^{m/2})$
 - **doubles** depth of search

Why is it called α - β ?

- α is the value of the best (i.e., highest-value) choice found so far at any choice point along the path for *max*
- If v is worse than α , *max* will avoid it
- \rightarrow prune that branch
- Define β similarly for *min*



The α - β algorithm

function ALPHA-BETA-SEARCH(*state*) *returns an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the *action* in SUCCESSORS(*state*) with value v

function MAX-VALUE(*state*, α , β) *returns a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for a, s in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** v

$\alpha \leftarrow \text{MAX}(\alpha, v)$

return v

The α - β algorithm

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  inputs: state, current state in game
            $\alpha$ , the value of the best alternative for MAX along the path to state
            $\beta$ , the value of the best alternative for MIN along the path to state

  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow +\infty$ 
  for  $a, s$  in SUCCESSORS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```

Evaluation functions

- For chess, typically **linear** weighted sum of **features**

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

e.g., $w_1 = 9$ with

$f_1(s) = (\text{number of white queens}) - (\text{number of black queens}), \text{ etc.}$

Cutting off search

MinimaxCutoff is identical to *MinimaxValue* except

1. *Terminal?* is replaced by *Cutoff?*
2. *Utility* is replaced by *Eval*

Does it work in practice?

$$b^m = 10^6, b=35 \rightarrow m=4$$

4-ply lookahead is a hopeless chess player!

- 4-ply \approx human novice
- 8-ply \approx typical PC, human master
- 12-ply \approx Deep Blue, Kasparov